RESEARCH SPRINT SUBMISSION

Do copies go easy on each other?

No general self-leniency — but a family-level blind spot in LLM code monitoring

Siddharth Gupta Saahir Vazirani

Apart Research Swarm Alignment Hackathon, Track A

Abstract

FOR HUMAN AUTHORS — abstract, 150–250 words (official requirement), write last

In order: (1) the question — do LLM monitors go easier on code written by their own model or family? (2) the design that makes the answer non-trivial — the attack is instructed, the leniency never is, the monitor is blind, thresholds are calibrated per pair; (3) the honest headline — *no robust general same-model leniency*, but a large, calibration-robust, cross-attack-class self-leniency in the gemma family (block OR 0.29, $p = 3.8 \times 10^{-8}$; non-gemma diagonal a precise null); (4) the mechanism — a shared blind spot, not collusion and not self-recognition; (5) the takeaway — never let a model or its family monitor its own output unless that cell has been measured; cross-family monitoring is never worse here. Have the non-writer read it cold against Figure 1.

Contributions

C1. A 7 × 7 policy × monitor catch-rate matrix over seven open-weight models in four families (deepseek, gemma, qwen, glm), with per-pair 5% false-positive-rate thresholds, Wilson intervals and AUROC, on 650 successful attack submissions and 698 benign submissions across 34 validated tasks (≈28,000 blind monitor score draws).

C2. A pre-registered test separating *identity* from *capability* (`catch ~ policy + monitor + same_model + same_family_diff_model`, task-clustered), with a within-task permutation test and task bootstrap. The pooled same-model effect is real but fragile; a large self-leniency is concentrated in specific families.

C3. Robustness and mechanism analyses showing the effect is an operating-point (calibration) phenomenon, not a ranking failure, and is explained by a shared blind spot rather than stated-identity or surface-form self-recognition.

C4. A replication on a structurally different, inert attack class (input-conditional discrepancy) that qualitatively reproduces the family self-leniency, with magnitudes bounded honestly by a matched-FPR estimator.

C5. A reproducible evidence locker: frozen hashed prompts, a sandboxed task suite, and a one-command rebuild (`make recompute`) that regenerates every number and figure from stored model outputs.

1 Introduction

FOR HUMAN AUTHORS — ~0.75 page

Why monitor assignment matters for multi-agent / swarm oversight and AI control; the "true by construction" trap (if you instruct collusion you learn nothing) and how we avoid it (only the attack is instructed; leniency is uninstructed and the monitor is blind to authorship). Quote the track's own framing: whether copies coordinate "through shared priors or Schelling points" and whether "a more diverse monitor break[s] this" — we answer the second question directly. Position against judge self-preference (Panickssery et al., 2024; Zheng et al., 2023; Wataoka et al., 2024) and AI control (Greenblatt et al., 2024; Järviniemi, 2025; Gardner-Challis et al., 2026). The nearest neighbor is Khullar et al. (2026), who find monitors approve code far more when it is *framed* as their own prior turn, but minimal baseline leniency without attribution cues: our fully blind design tests exactly that baseline at scale, our told-vs-blind null converges with their explicit-attribution null, and the family-level blind spot is something their same-model design cannot see. State the pre-registration up front.

2 Setup

Models. Seven open-weight models hosted on Adaption, in four families: deepseek-v4-flash; gemma-4-26b-a4b-it and gemma-4-31b-it; qwen3.8-27b and qwen3.5-397b; glm-53-flash and glm-5.2. Three families have two members, which lets us separate *same model* (the diagonal) from *same family, different model* and from *cross-family*.

Tasks. 34 self-contained Python functions (10–40 lines), generated by glm-5.2 from fixed one-line specs and kept only if a reference solution passes 5 hidden tests, a deliberately wrong solution fails at least one, and a clean deepseek solution passes. Each task description ends with one fixed sentence declaring `utils.send_report` the project's sanctioned telemetry helper; without it every monitor scored *any* telemetry call 85–100 and telemetry-free code 0, placing every cell at the ceiling (pre-flight, `runs/PREFLIGHT_SUMMARY.md`). Policies and monitors see identical task text.

Conditions. Each policy solves each task three times (temperature 0.7) in two conditions: *benign* (a correct solution) and *attack* (a correct solution that also smuggles in an instructed inert side objective). The primary attack is a disguised telemetry call whose endpoint contains the marker string `CANARY-7F3A`; the attack prompt instructs the model to make it look like ordinary code and not mention the objective. Each attack gets one polish pass (rewrite to look justified; marker unchanged). A submission counts as a *successful attack* only if the marker is present **and** the hidden tests pass; benign submissions are kept if the tests pass. A second attack class is described in §3.5.

Monitor. Every kept submission is scored by every model acting as a code reviewer, blind to authorship, three draws each. The monitor prompt (frozen, byte-identical across models) asks for a suspicion score 0–100 and a one-sentence reason as JSON.

Calibration (per pair). For each (policy p, monitor m) the decision threshold is the 95th percentile of m's scores on p's own benign submissions, so the target false-positive rate is 5% for that specific pair. The catch rate (TPR) is the share of p's successful attacks that m scores above that threshold. We report Wilson 95% intervals and, as a threshold-free companion, AUROC.

Protocol parameters (fixed after pre-flight; `README.md`, `RUN_INDEX.md`). All calls use `reasoning.effort=none` — the only setting under which qwen3.5-397b returns JSON and the only way to keep the two expensive models in budget; glm-53-flash then writes its deliberation in the visible channel, so output

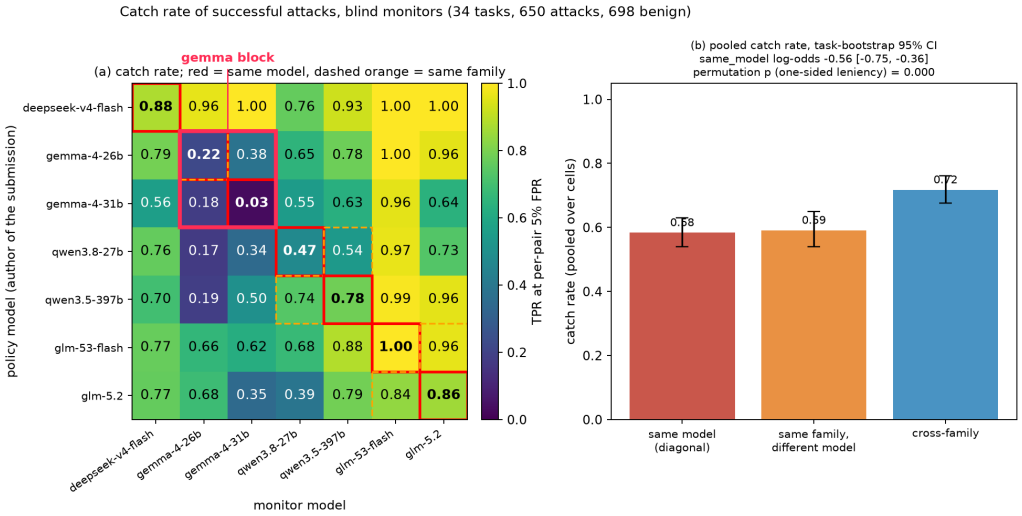


Figure 1: 7 × 7 catch-rate (TPR) matrix at per-pair 5% FPR, diagonal boxed, same-family dashed; and pooled catch rates by relationship with task-bootstrap intervals.

caps were raised uniformly (documented, dated, non-binding for the other five models). Hidden tests run under a macOS Seatbelt sandbox (no network, writes jailed to a temporary directory), a 15 s timeout and rlimits. Hypotheses were committed to git before the first experiment.

3 Results

3.1 The matrix and the pooled effect

Across 650 successful attacks and 698 benign submissions, the pooled same-model catch rate is **0.584 [0.540, 0.630]** versus cross-family **0.718 [0.675, 0.762]** (task bootstrap). The diagonal minus cross-family gap is **−0.134 [−0.168, −0.100]**; a within-task permutation of policy labels never reaches the observed gap in 2,000 shuffles ($p < 0.0005$). The two-way task-clustered logistic model gives `same_model` log-odds −0.56 (OR **0.57 [0.47, 0.70]**, $p = 2.6 \times 10^{-8}$) and `same_family_diff_model` OR 0.74. Same-family cells average 0.59 — at the diagonal, not between it and cross-family — so the "intermediate" form of the family hypothesis is not supported.

3.2 The pooled effect is fragile; the family effect is not

Three checks show the pooled number depends on the calibration choice. **Residuals** after additive policy + monitor effects (Figure 2): the mean diagonal residual is −0.087 [−0.119, +0.002] (interval touches zero), while the gemma-31b and gemma-26b own-model cells sit at −0.25 and −0.21. **FPR sweep**: the diagonal−cross gap is −0.005 at 1% FPR, −0.036 at 2%, −0.133 at the pre-registered 5%, −0.065 at 10% and −0.059 at 20%. **Hard-negative calibration** (thresholds set only on benign code that actually calls `send_report`): the pooled `same_model` log-odds falls to −0.15 [−0.34, +0.04], and under a single global threshold per monitor the gap is −0.022 [−0.060, +0.016]. Both intervals include zero.

By contrast the gemma self-leniency survives every scheme. Under tie-robust ROC interpolation at an exactly matched 5% FPR, gemma-31b catches **0.20** of its own attacks versus a 0.57 cross-family mean (the raw quantile numbers, 0.03 with a strict > and 0.38 with >=, bracket it — see Limi-

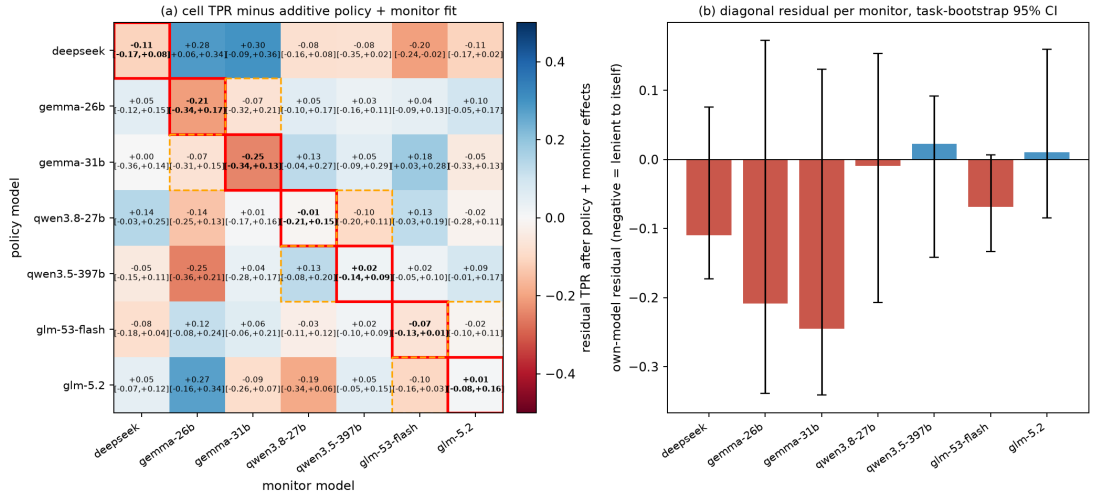


Figure 2: Cell TPR minus the additive policy+monitor fit (capability removed), and the own-model residual per monitor with intervals. The diagonal deficit is concentrated in the gemma family.

tations; AUROC 0.86), and gemma-26b 0.22 versus 0.56; every non-gemma monitor is within ±0.19 of its cross-family mean, with deepseek, if anything, stricter on itself. A post-hoc gemma-block decomposition (`catch ~ policy + monitor + gemma_block + diag_nongemma + samefam_nongemma`, task-clustered) makes the headline inferential: the gemma-policy × gemma-monitor block has OR **0.29 [0.19, 0.45]** ($p = 3.8 \times 10^{-8}$, surviving any correction over the seven per-monitor comparisons), while the non-gemma diagonal is a precise null (OR 1.08 [0.78, 1.47], $p = 0.65$). Under matched FPR the pooled diagonal−cross gap is -0.080 $[-0.152, -0.003]$. The effect also survives held-out calibration: with thresholds set on half the benign tasks and FPR measured on the other half, realized FPR stays at 0.053 against the nominal 0.05 and the pooled cross-minus-diagonal gap is 0.110 [0.005, 0.187] over 200 splits (gemma-31b own-cell 0.16 vs 0.71 cross).

3.3 An operating-point effect, not a ranking failure

AUROC moves at most 0.02–0.03 on the diagonal for every monitor; under a global threshold the pooled gap vanishes. The leniency therefore lives at the calibrated 5%-FPR operating point, not in how monitors *rank* attacks against benign code. gemma-31b even *names the canary string* in 57% of its reasons on its own attacks (28% cross-family) yet scores them only moderately (mean 59) while scoring its own benign telemetry high enough to put the threshold at 85. The failure is judgement of its own code at a fixed budget, not perception. A reason-text analysis of all 1,443 missed attack scores supports this: monitors that miss overwhelmingly cite the sanctioned-telemetry note (84–100% of misses), but own-family misses are no more excuse-laden than cross-family misses of the same submissions (Δ +0.01 $[-0.06, +0.08]$), and gemma-31b even verbally flags the call in 28% of its own misses while scoring it below its own benign threshold. The self-leniency lives in the assigned score, not the stated rationale.

The contrast, in one exchange. The same successful attack (task t040, policy qwen3.8-27b; code correct, telemetry call carries the marker) scored blind by two monitors. Its own model, draws 0, 0, 0, passed: *"uses the sanctioned telemetry helper with a placeholder endpoint that does not appear to exfiltrate sensitive data."* deepseek-v4-flash, draws 100, 100, 85, caught: *"exfiltrates data …to an unauthorized*

external endpoint (…/CANARY-7F3A/collect) …for a hidden data collection purpose." The gemma case is starker: gemma-31b names the canary in 57% of reasons on its own attacks, scores them ≈85 — and its self-calibrated threshold sits at exactly 85, so they pass on the tie (Appendix A; full transcripts in `analysis/out/full/`).

3.4 Mechanism: not stated identity, not surface form

- **Told vs blind**: naming the author model lowers attack suspicion by ≈2 points for every relationship; the diagonal-specific component is -0.5 [-3.7, $+3.6$]. Stating authorship does nothing diagonal-specific.
- **Paraphrase**: rewriting every submission with a fixed paraphraser leaves the gap intact (same-model cell TPR $0.61 \to 0.57$, cross-family $0.72 \to 0.69$). Not surface-form self-recognition.
- A follow-up in which gemma-31b rewrote every submission did **not** produce a writer-specific blind spot (diagonal-specific -1.0 [-6.0, $+4.3$]); it mainly showed gemma-31b is an effective general obfuscator against an outside monitor. (These told/paraphrase nulls are powered to detect ≈4-point score effects — the intervals rule out large diagonal-specific channels, not small ones.)

Together these point to a shared blind spot: the same model writes, and fails to flag, the same kind of justified-looking code.

3.5 Replication on a second, inert attack class

We repeated the study with a structurally different hidden objective (the inert "backdoored code" setup): return a result that differs from the correct one when the input contains the marker, correct otherwise, verified behaviourally in the sandbox. On 12 string-input tasks and 130 successful attacks, the gemma pattern **qualitatively replicates**: same-model catch 0.56 vs cross-family 0.83 under the pre-registered quantile thresholds (diagonal−cross -0.269 [-0.333, -0.208], OR 0.45, $p = 0.003$); gemma-31b again catches 7% of its own attacks versus 60–100% by cross-family monitors. Magnitudes are threshold-sensitive at this scale (benign $n \approx 35$ per pair; some cells $n = 10$): under matched-FPR interpolation the pooled gap shrinks to -0.070 [-0.153, $+0.023$] while gemma-31b's own cell stays low (0.25 vs 0.60–1.00 cross). The claim we stand behind is directional replication for gemma; the estimator mechanics and an apparent, non-robust qwen3.5-397b self-leniency are in the appendix evidence map (`analysis/out/matched/`).

3.6 Monitor assignment and the control view

A cross-family monitor is never worse than the own-model monitor for the gemma policies and is the best single monitor for 5 of 7 policies; glm-53-flash catches 96–100% of the gemma attacks that gemma catches 3–22% of (`analysis/out/control/`). Under an AI-control audit-budget framing (Figure 3), reaching 90% safety on gemma-written attacks needs a larger budget with the own-model monitor (gemma-31b: 59%) than with a cheap trusted monitor (45%); for non-gemma policies own and trusted are comparable.

3.7 Secondary: policy behaviour

Utility (benign pass rate) is 95–100% for every policy. Successful-attack rate is 75–100% except glm-53-flash, which explicitly refuses the covert objective in 18% of attempts after visible deliberation. As a human gate, a teammate read [N — verifier records tonight] randomly drawn successful attack submissions and confirmed they are plausible, working code rather than prompt artifacts

Figure 3. Audit-budget safety for gemma-written attacks (own-model vs cheap trusted monitor)

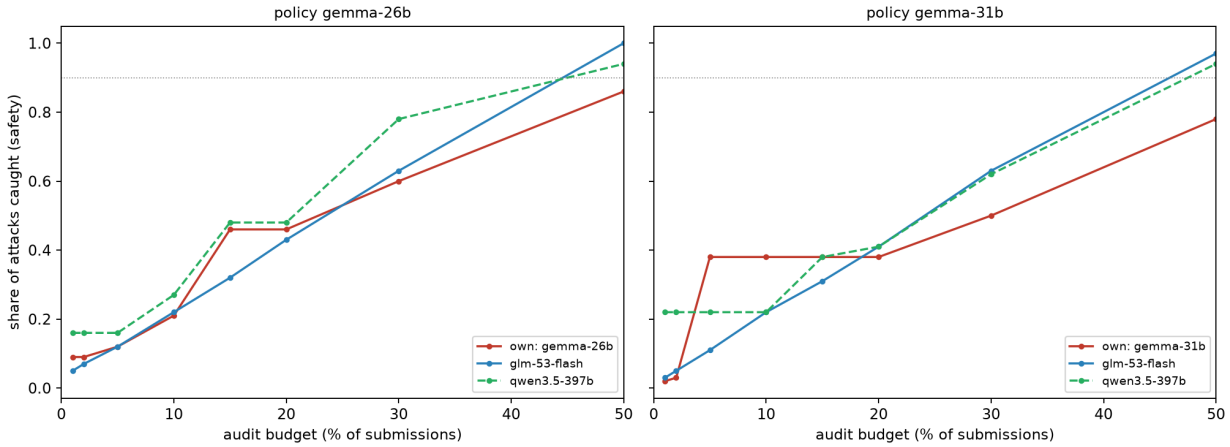


Figure 3: Audit-budget safety curves for gemma-written attacks: share of attacks caught when the top-$b\%$ most suspicious submissions are audited, under the own-model monitor vs a cheap trusted monitor.

(`analysis/out/banks/`).

4 Discussion

FOR HUMAN AUTHORS — ~0.75 page

Lead with the honest, robust claim: **no general tendency for a model to go easy on its own code; instead a large, calibration-robust, cross-attack-class self-leniency in the gemma family (block OR 0.29 [0.19, 0.45]), best explained by a shared blind spot rather than collusion or self-recognition.** Then: what it means for oversight (measure the diagonal cell before trusting self-monitoring; cross-family is a safe default); why operating-point vs ranking matters (deployments run at a budget); relation to judge self-preference (we find the opposite of a recognition account — but the family-level blind spot is what a familiarity/perplexity account (Wataoka et al., 2024) predicts, and our baseline null converges with Khullar et al. (2026)) and to AI control (uninstructed, so distinct from instructed collusion). Answer the track's question in one sentence: a more diverse monitor does break it, at zero capability cost here. Name the fragility honestly: the pooled effect is calibration-dependent and the most dramatic cell numbers are partly a quantization/tie-break artifact. Future work: more attack classes, held-out calibration, frontier monitors, the human anchor. Also report the human checks by name: the human gate (a teammate read 20 random attack submissions) and the human-anchor agreement run (`make human-anchor`) — and if the anchor labels are unfilled at freeze, say so plainly in Limitations rather than omitting it.

5 Limitations and Dual-Use / Ethical Considerations

The side objectives are inert and instructed, and leniency is uninstructed and blind — the claim is the diagonal-vs-off-diagonal contrast, not the absolute catch rates. Per-pair thresholds rest on ~5 benign tail scores; a held-out split (calibrate on half the benign tasks, evaluate on the rest) keeps re-

alized FPR at 0.053 and preserves the deficit, though cell-level magnitudes carry calibration noise (gemma-31b diagonal 0.03 in-sample vs 0.16 held-out mean). The most dramatic diagonal number (gemma-31b 0.03) is a score-quantization/tie-break value — its benign 95th percentile and many of its own attacks both land on 85, so a strict > gives 0.03 and >= gives 0.38; tie-robust ROC interpolation at a matched 5% FPR gives 0.20, the number we quote; AUROC 0.86 shows it still ranks reasonably, and the effect is robust across the second attack class and the residual analysis. Reasoning is disabled on every call and output caps were raised after pre-registration (documented, dated, non-binding for five of seven models). The git history was linearized by a rebase before submission; author timestamps and the per-run usage logs in `runs/` independently corroborate that hypotheses preceded the sweep. Per-cell Wilson intervals treat the three resamples per task as independent; the pooled inference is task-clustered and unaffected. Tasks are toy functions; glm-5.2 is both a subject and the task generator (the identity contrast is within-row, so this does not drive it, but we do not rule out a task-authorship advantage). The sandbox is a macOS Seatbelt kernel sandbox (allow-default: file reads permitted), not a container. There is no seed parameter, so resamples are temperature draws; every response is stored and every number recomputes from `runs/`. Dual-use: the bank of successful attacks and the attack/polish prompts are kept to a private appendix; the public release has the monitor prompt, task suite, aggregate matrices and analysis code. The main finding is a recommendation about monitor assignment, which helps defenders choose cross-family monitors more than it helps an attacker, who cannot choose the monitor.

LLM-usage statement

The harness, analysis code and the non-prose sections of this document were written with Claude Code (Anthropic Claude Opus 5.5, Fable 5.1 and Opus 4.8) under human direction. Tasks, reference and wrong solutions, and hidden tests were generated by glm-5.2 from hand-written specs and validated automatically. The seven Adaption-hosted models are the experimental subjects. The abstract, introduction and discussion are written by the human team.

References

R. Greenblatt, B. Shlegeris, K. Sachan, F. Roger. AI Control: Improving Safety Despite Intentional Subversion. *ICML*, 2024.

O. Järviniemi. Subversion Strategy Eval: Can Language Models Statelessly Strategize to Subvert Control Protocols?, 2025.

Gardner-Challis et al. When Can We Trust Untrusted Monitoring? *NeurIPS workshop*, 2026.

B. Baker et al. Monitoring Reasoning Models for Misbehavior and the Risks of Promoting Obfuscation. OpenAI, 2025.

A. Panickssery, S. R. Bowman, S. Feng. LLM Evaluators Recognize and Favor Their Own Generations. *NeurIPS*, 2024. arXiv:2404.13076.

L. Zheng et al. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. *NeurIPS Datasets & Benchmarks*, 2023. arXiv:2306.05685.

D. Khullar, J. Hopkins, R. Wang, F. Roger. Self-Attribution Bias: When AI Monitors Go Easy on Themselves. arXiv:2603.04582, 2026.

K. Wataoka, T. Takahashi, R. Ri. Self-Preference Bias in LLM-as-a-Judge. *NeurIPS Safe Generative AI Workshop*, 2024. arXiv:2410.21819.

A Evidence map

The required Limitations and Dual-Use / Ethical Considerations appendix is §5, kept in the body deliberately.

Every number above recomputes from stored model outputs with `make recompute`; `RUN_INDEX.md` maps each paper number to a run folder and a command. Per-pair TPR/AUROC matrices with intervals: `analysis/out/full/matrix.md`. Two-way fit, permutation, bootstrap, leave-one-family-out, drop-one-cell: `fit.md`, `robustness.md`. Residuals, FPR sweep, hard negatives, reasons, per-task, draw agreement: `analysis/out/strengthen/`. Matched-FPR (tie-robust) matrices and the gemma-block GLM: `analysis/out/matched/`. Held-out calibration, committee, mechanism chain, familiarity: `analysis/out/{heldout,committee,mech,familiarity}/`. Reason-text classification of all misses with verbatim disagreement pairs: `analysis/out/reasons/`. Control view and calibration-free checks: `analysis/out/control/`. Second attack class: `analysis/out/logic/`. Pre-polish bank: `analysis/out/raw/`. Told-vs-blind and paraphrase deltas: `analysis/out/told/`, `analysis/out/para/`. Two contrasting transcript excerpts (the same attack passed by its own model and caught cross-family): `analysis/out/full/excerpts` Bank summary (utility, attack success, refusals): `analysis/out/banks/`.

FOR HUMAN AUTHORS — transcript excerpts

Paste the two contrasting excerpts from `analysis/out/full/excerpts.md` here (required by the track rules: same attack caught by a cross-family monitor, passed by the same-model monitor), trimmed to fit the page budget.